

Overview

- 6.828 goals
 - Understand operating systems in detail by designing and implementing a small O/S
 - Hands-on experience with building systems ("Applying 6.033")
- What do applications want from an O/S?
 - Abstract the hardware for convenience and portability
 - Multiplex the hardware among multiple applications
 - Isolate applications to contain bugs
 - Allow sharing among applications
- What is an OS?
 - e.g. OSX, Windows, Linux
 - the small view: a h/w management library
 - the big view: physical machine -> abstract one w/ better properties
- Organization: layered picture h/w: CPU, mem, disk kernel: [various services] user: applications, e.g. vi and gcc
 - we care a lot about the interfaces and internal kernel structure
- What services does an O/S kernel typically provide?
 - processes
 - memory
 - file contents
 - directories and file names
 - security
 - many others: users, IPC, network, time, terminals
- What does an O/S abstraction look like?
 - Applications only see them via system calls
 - Examples, from UNIX / Linux:

```
fd = open("out", 1);
write(fd, "hello\n", 6);
pid = fork();
```
- Why is O/S design/implementation hard/interesting?
 - the environment is unforgiving: weird h/w, no debugger
 - it must be efficient (thus low-level?) ...but abstract/portable (thus high-level?)
 - powerful (thus many features?) ...but simple (thus a few composable building blocks?)
 - features interact: `fd = open(); ...; fork()`
 - behaviors interact: CPU priority vs memory allocator.
 - open problems: security, multi-core
- You'll be glad you learned about operating systems if you...
 - want to work on the above problems
 - care about what's going on under the hood

- have to build high-performance systems
- need to diagnose bugs or security problems

Class structure

- See web site: <http://pdos.lcs.mit.edu/6.828>
- Lectures
 - basic O/S ideas
 - extended inspection of xv6, a traditional O/S
 - several more recent topics
 - xv6 programming to re-inforce xv6 understanding
- Lab: JOS, a small O/S for x86 in an exokernel style
 - you build it, 5 labs + final lab of your choice
 - kernel interface: expose hardware, but protect -- no abstractions!
 - unprivileged library: fork, exec, pipe, ...
 - applications: file system, shell, ..
 - development environment: gcc, qemu
 - lab 1 is out
 - make grade
- Code review
- Two quizzes: one in class hours, one in final's week

Shell and system calls

- 6.828 is largely about design and implementation of system call interface. let's start by looking at how programs use that interface. example: the Unix shell.
- the shell is the Unix command UI
 - typically handles login session, runs other processes
 - you saw it in 6.033: <http://web.mit.edu/6.033/www/assignments/handson-unix.html>
 - the shell is also a programming/scripting language
 - look at some simple examples of shell operations, how they use different O/S abstractions, and how those abstractions fit together. See [Unix paper](#) if you are unfamiliar with the shell.
- [Simplified xv6 sh.c](#)
 - See [chapter 0 of xv6 book](#)
 - Basic organization: parsing and executing commands (e.g., ls, ls | wc, ls > out)
 - Shell implemented using system calls (e.g., read, write, fork, exec, wait) conventions: -1 return value signals error, error code stored in `errno`, `perror` prints out a descriptive error message based on `errno`.
 - Many systems calls are encapsulated in libc calls (e.g., `fgets` vs `read`)
- Trace system calls `$ ls`
 - On OSX: `sudo dtruss ./a.out` (where `a.out` is the compiled `sh.c`)
 - On Linux: `strace ./a.out`

- what does fork() do? copies user memory copies process kernel state (e.g. user id) child gets a different PID child state contains parent PID returns twice, with different values
- parent and child may run concurrently (e.g., on different processors).
- what does wait() do? waits for any child to exit what if child exits before parent calls wait?
- what are file descriptors? (0, 1, 2, etc. in read/write) [echo.c](#)
- what is i/o redirection? [echo.c](#) How would you implement ">" in sh.c
- what are pipes? (ls | wc) [pipe1.c](#) [pipe2.c](#) How would you implement them in sh.c?
- Homework assignment for [shell](#) Edit By [MaHua](#)