

# 6.828 Lecture Notes: x86 and PC architecture

## Outline

- PC architecture
- x86 instruction set
- gcc calling conventions
- PC emulation

## PC architecture

- A full PC has:
  - an x86 CPU with registers, execution unit, and memory management
  - CPU chip pins include address and data signals
  - memory
  - disk
  - keyboard
  - display
  - other resources: BIOS ROM, clock, ...
- We will start with the original 16-bit 8086 CPU (1978)
- CPU runs instructions:

```
for(;;){  
    run next instruction  
}
```

- Needs work space: registers
  - four 16-bit data registers: AX, BX, CX, DX
  - each in two 8-bit halves, e.g. AH and AL
  - very fast, very few
- More work space: memory
  - CPU sends out address on address lines (wires, one bit per wire)
  - Data comes back on data lines
  - or data is written to data lines
- Add address registers: pointers into memory
  - SP - stack pointer
  - BP - frame base pointer
  - SI - source index
  - DI - destination index
- Instructions are in memory too!
  - IP - instruction pointer (PC on PDP-11, everything else)
  - increment after running each instruction
  - can be modified by CALL, RET, JMP, conditional jumps
- Want conditional jumps
  - FLAGS - various condition codes
    - whether last arithmetic operation overflowed
    - ... was positive/negative
    - ... was [not] zero
    - ... carry/borrow on add/subtract

- ... etc.
- whether interrupts are enabled
- direction of data copy instructions
- JP, JN, J[N]Z, J[N]C, J[N]O ...
- Still not interesting - need I/O to interact with outside world
  - Original PC architecture: use dedicated *I/O space*
    - Works same as memory accesses but set I/O signal
    - Only 1024 I/O addresses
    - Accessed with special instructions (IN, OUT)
    - Example: write a byte to line printer:

```
#define DATA_PORT    0x378
#define STATUS_PORT   0x379
#define    BUSY 0x80
#define CONTROL_PORT 0x37A
#define    STROBE 0x01
void
lpt_putc(int c)
{
    /* wait for printer to consume previous byte */
    while((inb(STATUS_PORT) & BUSY) == 0)
        ;

    /* put the byte on the parallel lines */
    outb(DATA_PORT, c);

    /* tell the printer to look at the data */
    outb(CONTROL_PORT, STROBE);
    outb(CONTROL_PORT, 0);
}
```

### Memory-Mapped I/O

Use normal physical memory addresses

Gets around limited size of I/O address space  
 No need for special instructions  
 System controller routes to appropriate device

Works like ``magic'' memory:

*Addressed and accessed like memory,*  
*but ...*  
*... does not behave like memory!*  
*Reads and writes can have ``side effects''*  
*Read results can change due to external*

events

What if we want to use more than  $2^{16}$  bytes of memory?

8086 has 20-bit physical addresses, can have 1 Meg RAM  
 the extra four bits usually come from a 16-bit "segment  
 register":

CS - code segment, for fetches via IP  
 SS - stack segment, for load/store via SP and BP  
 DS - data segment, for load/store via other registers  
 ES - another data segment, destination for string

operations

virtual to physical translation:  $pa = va + seg * 16$

e.g. set CS = 4096 to execute starting at 65536  
 tricky: can't use the 16-bit address of a stack variable as  
 a pointer  
 a *far pointer* includes full segment:offset (16 + 16 bits)  
 tricky: pointer arithmetic and array indexing across  
 segment boundaries

But 8086's 16-bit addresses and data were still painfully small

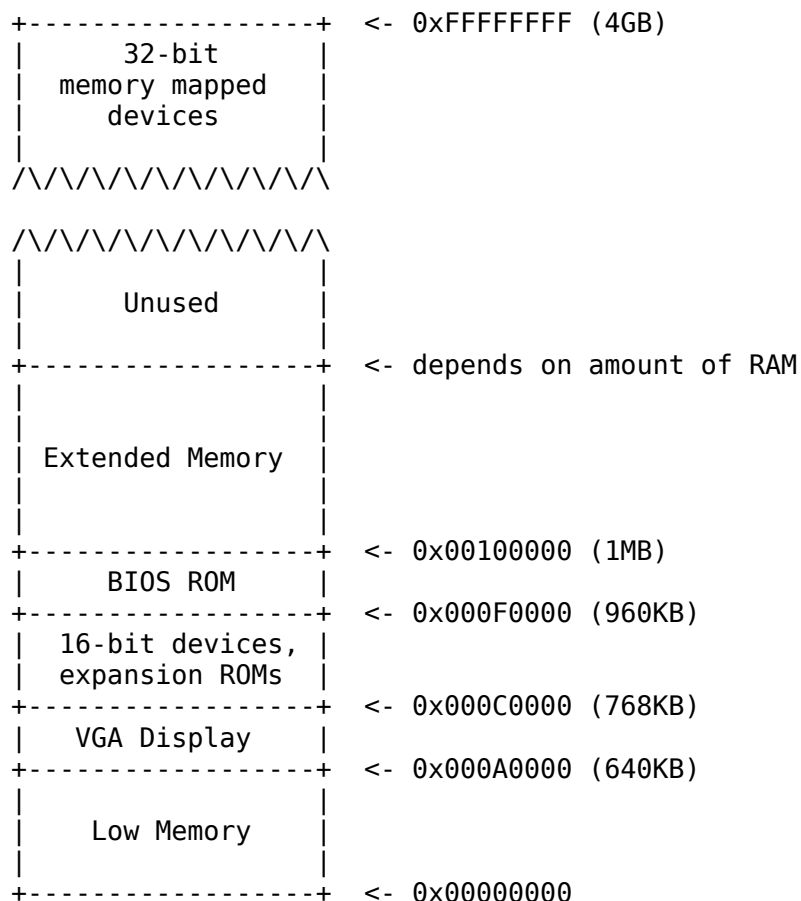
80386 added support for 32-bit data and addresses (1985)  
 boots in 16-bit mode, boot.S switches to 32-bit mode  
 registers are 32 bits wide, called EAX rather than AX  
 operands and addresses that were 16-bit became 32-bit in 32-bit  
 mode, e.g. ADD does 32-bit arithmetic  
 prefixes 0x66/0x67 toggle between 16-bit and 32-bit operands and  
 addresses: in 32-bit mode, MOVW is expressed as 0x66 MOVW  
 the .code32 in boot.S tells assembler to generate 0x66 for e.g.  
 MOVW  
 80386 also changed segments and added paged memory...

Example instruction encoding

b8 cd ab	16-bit CPU, AX <- 0xabcd
b8 34 12 cd ab	32-bit CPU, EAX <- 0xabcd1234
66 b8 cd ab	32-bit CPU, AX <- 0xabcd

## x86 Physical Memory Map

- The physical address space mostly looks like ordinary RAM
- Except some low-memory addresses actually refer to other things
- Writes to VGA memory appear on the screen
- Reset or power-on jumps to ROM at 0xfffff0 (so must be ROM at top...)



## x86 Instruction Set

- Intel syntax: op dst, src (Intel manuals!)
- AT&T (gcc/gas) syntax: op src, dst (labs, xv6)
  - uses b, w, l suffix on instructions to specify size of operands
- Operands are registers, constant, memory via register, memory via constant
- Examples:

<u>AT&amp;T syntax</u>	<u>"C"-ish equivalent</u>	
movl %eax, %edx	edx = eax;	<i>register mode</i>
movl \$0x123, %edx	edx = 0x123;	<i>immediate</i>
movl 0x123, %edx	edx = *(int32_t*)0x123;	<i>direct</i>
movl (%ebx), %edx	edx = *(int32_t*)ebx;	<i>indirect</i>
movl 4(%ebx), %edx	edx = *(int32_t*)(ebx+4);	<i>displaced</i>

- Instruction classes
  - data movement: MOV, PUSH, POP, ...
  - arithmetic: TEST, SHL, ADD, AND, ...
  - i/o: IN, OUT, ...
  - control: JMP, JZ, JNZ, CALL, RET
  - string: REP MOVSB, ...
  - system: IRET, INT
- Intel architecture manual Volume 2 is *the* reference

## gcc x86 calling conventions

- x86 dictates that stack grows down:

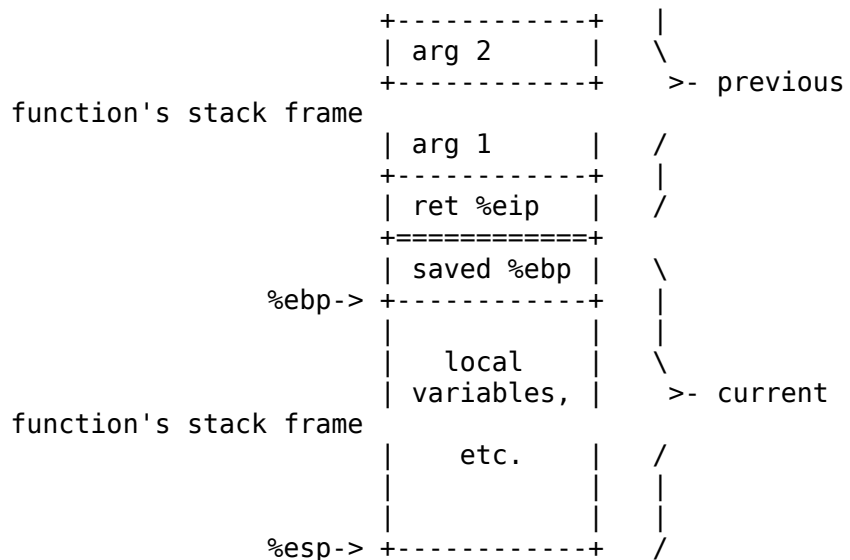
Example instruction    What it does

pushl %eax	subl \$4, %esp movl %eax, (%esp)
popl %eax	movl (%esp), %eax addl \$4, %esp
call 0x12345	pushl %eip (*) movl \$0x12345, %eip (*)
ret	popl %eip (*)

(\*) *Not real instructions*

- GCC dictates how the stack is used. Contract between caller and callee on x86:
  - at entry to a function (i.e. just after call):
    - %eip points at first instruction of function
    - %esp+4 points at first argument
    - %esp points at return address
  - after ret instruction:
    - %eip contains return address

- %esp points at arguments pushed by caller
- called function may have trashed arguments
- %eax (and %edx, if return type is 64-bit) contains return value (or trash if function is void)
- %eax, %edx (above), and %ecx may be trashed
- %ebp, %ebx, %esi, %edi must contain contents from time of call
- Terminology:
  - %eax, %ecx, %edx are "caller save" registers
  - %ebp, %ebx, %esi, %edi are "callee save" registers
- Functions can do anything that doesn't violate contract. By convention, GCC does more:
  - each function has a stack frame marked by %ebp, %esp



- %esp can move to make stack frame bigger, smaller
- %ebp points at saved %ebp from previous function, chain to walk stack
- function prologue:

```

pushl %ebp
movl %esp, %ebp

```

or

```

enter $0, $0

```

enter usually not used: 4 bytes vs 3 for pushl+movl, not on hardware fast-path anymore

- function epilogue can easily find return EIP on stack:

```

movl %ebp, %esp
popl %ebp

```

or

```

leave

```

leave used often because it's 1 byte, vs 3 for movl+popl

- Big example:

- C code

```
int main(void) { return f(8)+1; }
int f(int x) { return g(x); }
int g(int x) { return x+3; }
```

- assembler

```
_main:
    pushl %ebp                prologue
    movl %esp, %ebp
    pushl $8                  body
    call _f
    addl $1, %eax
    movl %ebp, %esp           epilogue
    popl %ebp
    ret

_f:
    pushl %ebp                prologue
    movl %esp, %ebp
    pushl 8(%esp)             body
    call _g
    movl %ebp, %esp           epilogue
    popl %ebp
    ret

_g:
    pushl %ebp                prologue
    movl %esp, %ebp
    pushl %ebx                save %ebx
    movl 8(%ebp), %ebx        body
    addl $3, %ebx
    movl %ebx, %eax
    popl %ebx                 restore %ebx
    movl %ebp, %esp           epilogue
    popl %ebp
    ret
```

- Super-small \_g:

```
_g:
    movl 4(%esp), %eax
    addl $3, %eax
    ret
```

- Shortest \_f?
- Compiling, linking, loading:
  - *Preprocessor* takes C source code (ASCII text), expands `#include` etc, produces C source code
  - *Compiler* takes C source code (ASCII text), produces assembly language (also ASCII text)

- *Assembler* takes assembly language (ASCII text), produces .o file (binary, machine-readable!)
- *Linker* takes multiple '.o's, produces a single *program image* (binary)
- *Loader* loads the program image into memory at run-time and starts it executing

## PC emulation

- The Bochs emulator works by
  - doing exactly what a real PC would do,
  - only implemented in software rather than hardware!
- Runs as a normal process in a "host" operating system (e.g., Linux)
- Uses normal process storage to hold emulated hardware state: e.g.,
  - Stores emulated CPU registers in global variables

```
int32_t regs[8];
#define REG_EAX 1;
#define REG_EBX 2;
#define REG_ECX 3;
...
int32_t eip;
int16_t segregs[4];
...
```

- Stores emulated physical memory in Bochs's memory

```
char mem[256*1024*1024];
```

- Execute instructions by simulating them in a loop:

```
for (;;) {
    read_instruction();
    switch (decode_instruction_opcode()) {
    case OPCODE_ADD:
        int src = decode_src_reg();
        int dst = decode_dst_reg();
        regs[dst] = regs[dst] + regs[src];
        break;
    case OPCODE_SUB:
        int src = decode_src_reg();
        int dst = decode_dst_reg();
        regs[dst] = regs[dst] - regs[src];
        break;
    ...
    }
    eip += instruction_length;
}
```

- Simulate PC's physical memory map by decoding emulated "physical" addresses just like a PC would:

```
#define KB            1024
#define MB            1024*1024

#define LOW_MEMORY    640*KB
#define EXT_MEMORY    10*MB

uint8_t low_mem[LOW_MEMORY];
uint8_t ext_mem[EXT_MEMORY];
```

```

uint8_t bios_rom[64*KB];

uint8_t read_byte(uint32_t phys_addr) {
    if (phys_addr < LOW_MEMORY)
        return low_mem[phys_addr];
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        return rom_bios[phys_addr - 960*KB];
    else if (phys_addr >= 1*MB && phys_addr <
1*MB+EXT_MEMORY) {
        return ext_mem[phys_addr-1*MB];
    }
    else ...
}

void write_byte(uint32_t phys_addr, uint8_t val) {
    if (phys_addr < LOW_MEMORY)
        low_mem[phys_addr] = val;
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        ; /* ignore attempted write to ROM! */
    else if (phys_addr >= 1*MB && phys_addr <
1*MB+EXT_MEMORY) {
        ext_mem[phys_addr-1*MB] = val;
    }
    else ...
}

```

- Simulate I/O devices, etc., by detecting accesses to "special" memory and I/O space and emulating the correct behavior: e.g.,
  - Reads/writes to emulated hard disk transformed into reads/writes of a file on the host system
  - Writes to emulated VGA display hardware transformed into drawing into an X window
  - Reads from emulated PC keyboard transformed into reads from X input event queue