

Homework solution ==

- Let's review solution for [sh.c](#)
 - exec
 - redirect
 - pipe
- The process graph for a complicated pipeline
 - Who waits for whom? (draw a tree of processes)
 - Why close read-end and write-end? ensure that every process starts with 3 file descriptors, and that reading from the pipe returns end of file after the first command exits.

Exploring system calls with more shell features ==

- You can run the shell, redirect its stdin/stdout, etc.
 - I'll run this shell script with `sh < script`:

```
echo one
echo two
```
 - What will this shell command do? `$ sh < script > out`
 - the script itself didn't redirect the echo output, but it did inherit a fd 1 that was redirected to out.
 - to make this work it is important that offset is implicit (maintained by kernel)
 - this is why read and write don't take an offset argument
 - Is the following the same as above?

```
$ echo one > out
$ echo two > out
```
 - How to implement sequencing/lists (`;&&,||`)

```
$ gcc x.c ; ./a.out

$ gcc x.c && ./a.out
```
 - How to implement nesting?

```
$ (echo one; echo two) > out
```
- How does the shell implement "&"? `$ sleep 2 &`
 - Q: What if a background process exits while sh waits for a foreground process?

System call observations ==

- The fork/exec split looks wasteful; why is it useful? (A: exercise 2)

- System call interface simple, just ints and char buffers. why not have open() return a pointer reference to a kernel file object?
- Linux has a nice representation of a process and its FDs, under /proc/PID/
 - maps: VA range, perms (p=private, s=shared), offset, dev, inode, pathname
 - fd: symlinks to files pointed to by each fd.
- The file descriptor design:
 - nice interaction with fork
 - FDs help make programs more general purpose: don't need special cases for files vs console vs pipe
 - shell pipelines only work for programs w/ common formats (lines of text)

OS organization ==

- Now we have a feel for what Unix system call interface provides, how to implement the interface?
- Why have an o/s at all? why not just a library? then apps are free to use it, or not -- flexible apps can directly interact with hardware some tiny O/Ss for embedded processors work this way
- Key requirement: isolation multiplexing interaction
- helpful approach: abstract machine resources rather than raw hardware File system, not raw disk TCP, not a raw ethernet Processes, not raw CPU/memory abstractions often ease multiplexing and interaction also more convenient and portable
- Goals: apps must use OS interface, cannot directly interact with hardware apps cannot harm operating system
- Hardware support for isolation
 - Processors support user mode and kernel mode
 - some instructions can only be executed in kernel mode e.g., instructions to directly interact with hardware
 - If an application executes a privileged instruction hardware doesn't allow it
 - instead switches to kernel mode
 - kernel can clean up
- Leverage hardware support
 - Operating systems runs in kernel mode
 - kernel is a big program services: processes, file system, net low-level: devices, virtual memory all of kernel runs with full hardware privilege (convenient)
 - Applications run in user mode
 - isolated from kernel
 - systems calls switch between user and kernel mode application call instructions to enter kernel instruction enters kernel at an entry point specified by kernel
- What to run in kernel mode?
 - xv6 follows a traditional design: all of the OS runs in kernel mode
 - this design is called a monolithic kernel

- kernel interface \approx system call interface
- good: easy for subsystems to cooperate one cache shared by file system and virtual memory
- bad: interactions are complex leads to bugs no isolation within kernel
- alternative: microkernel design
 - many OS services run as ordinary user programs file system in a file server
 - kernel implements minimal mechanism to run services in user space IPC virtual memory threads
 - kernel interface \neq system call interface
 - good: more isolation
- jos: doesn't abstract hardware resources, but still provide isolation
- Can one have process isolation WITHOUT h/w-supported kernel/user mode? yes! see Singularity O/S, later in semester but h/w user/kernel mode is the most popular plan

xv6 kernel address space ==

- boundary between apps and kernel kernel/user bit kernel address space ...
- start w. kernel address space machine has booted code runs without virtual memory set up up kernel address space walk through kvmalloc()

Edit By [MaHua](#)