

6.828 2014 Lecture 6: Virtual Memory ==

- plan: address spaces paging hardware xv6 VM code
- today's problem: [user/kernel diagram] suppose the shell has a bug: sometimes it writes to a random memory address how can we keep it from wrecking the kernel? and from wrecking other processes?
- we want isolated address spaces each process has its own memory it can read and write its own memory it cannot read or write anything else
- xv6 uses x86's paging hardware to implement AS's ask questions! this material is important
- paging provides a level of indirection for addressing CPU -> MMU -> RAM

VA PA

s/w can only ld/st to virtual addresses, not physical kernel tells MMU how to map each virtual address to a physical address MMU essentially has a table, indexed by va, yielding pa called a "page table" MMU can restrict what virtual addresses user code can use

- x86 maps 4-KB "pages" and aligned -- start on 4 KB boundaries thus page table index is top 20 bits of VA
- what is in a page table entry (PTE)? see handout top 20 bits are top 20 bits of physical address "physical page number" MMU replaces top 20 of VA with PPN low 12 bits are flags Present, Writeable, &c
- where is the page table stored? in RAM -- MMU loads (and stores) PTEs o/s can read/write PTEs
- would it be reasonable for page table to just be an array of PTEs? how big is it? 2^{20} is a million 32 bits per entry 4 MB for a full page table -- pretty big on early machines would waste lots of memory for small programs! you only need mappings for a few hundred pages so the rest of the million entries would be there but not needed
- x86 uses a "two-level page table" to save space diagram pages of PTEs in RAM page directory (PD) in RAM PDE also contains 20-bit PPN -- of a page of 1024 PTEs 1024 PDEs point to PTE pages each PTE page has 1024 PTEs -- so 1024×1024 PTEs in total PD entries can be invalid those PTE pages need not exist so a page table for a small address space can be small
- how does the mmu know where the page table is located in RAM? %cr3 holds phys address of PD PD holds phys address of PTE pages they can be anywhere in RAM -- need not be contiguous
- how does x86 paging hardware translate a va? need to find the right PTE %cr3 points to PA of PD top 10 bits index PD to get PA of PT next 10 bits index PT to get PTE PPN from PTE + low-12 from VA
- flags in PTE P, W, U xv6 uses U to forbid user from using kernel memory
- what if P bit not set? or store and W bit not set? "page fault" CPU saves registers, forces transfer to kernel trap.c in xv6 source kernel can just produce error, kill process or kernel can install a PTE, resume the process e.g. after loading the page of memory from disk
- Q: why mapping rather than e.g. base/bound? indirection allows paging h/w to solve many problems e.g. avoids fragmentation e.g. copy-on-write fork e.g. lazy allocation (home work for

next lecture) many more techniques topic of next lecture

- how does xv6 use the x86 paging hardware?
- big picture of an xv6 address space -- one per process [diagram] 0x00000000:0x80000000 -- user addresses below KERNBASE 0x80000000:0x80100000 -- map low 1MB devices (for kernel) 0x80100000:? -- kernel instructions/data ? :0x8E000000 -- 224 MB of DRAM mapped here 0xFE000000:0x00000000 -- more memory-mapped devices
- where does xv6 map these regions, in phys mem? [diagram] note double-mapping of user pages
- each process has its own address space and its own page table all processes have the same kernel (high memory) mappings kernel switches page tables (i.e. sets %cr3) when switching processes
- Q: why this address space arrangement? user virtual addresses start at zero of course user va 0 maps to different pa for each process 2GB for user heap to grow contiguously but needn't have contiguous phys mem -- no fragmentation problem both kernel and user mapped -- easy to switch for syscall, interrupt kernel mapped at same place for all processes eases switching between processes easy for kernel to r/w user memory using user addresses, e.g. sys call arguments easy for kernel to r/w physical memory pa x mapped at va x+0x80000000 we'll see this soon while manipulating page tables
- Q: what's the largest process this scheme can accomodate?
- Q: could we increase that by increasing/decreasing 0x80000000?
- To think about: does the kernel have to map all of phys mem into its virtual address space?
- let's look at some xv6 virtual memory code virtual memory == address space / translation

** a process calls sbrk(n) to ask for n more bytes of heap memory malloc() uses sbrk() each process has a size kernel adds new memory at process's end, increases size sbrk() allocates physical memory (RAM) maps it into the process's page table returns the starting address of the new memory

** sys_sbrk() in sysproc.c

** growproc() in proc.c proc->sz is the process's current size allocvm() does most of the work switchvm sets %cr3 with new page table also flushes some MMU caches so it will see new PTEs

** allocvm() in vm.c why if(newsz >= KERNBASE) ? why PGROUNDUP? arguments to mappages()...

** mappages() in vm.c arguments are PD, va, size, pa, perm adds mappings from a range of va's to corresponding pa's rounds b/c some users pass in non-page-aligned addresses for each page-aligned address in the range call walkpgdir to find address of PTE need the PTE's address (not just content) b/c we want to modify put the desired pa into the PTE mark PTE as valid w/ PTE_P

** diagram of PD &c, as following steps build it

* walkpgdir() in vm.c mimics how the paging h/w finds the PTE for an address refer to the handout PDX extracts top ten bits &pgdir[PDX(va)] is the address of the relevant PDE now pde is the PDE if PTE_P the relevant page-table page already exists PTE_ADDR extracts the PPN from the PDE p2v() adds 0x80000000, since PTE holds physical address if not PTE_P alloc a page-table page fill in PDE with PPN -- thus v2p now the PTE we want is in the page-table page at offset PTX(va) which is 2nd 10 bits of va
Edit By [MaHua](#)