

plan

homework sequence coordination xv6: sleep & wakeup challenges lost wakeup problem signals

homework: context switching

user-level context switching same idea as kernel-level show solution Q: what address is 161 Q: why is it on the stack? Q: what happens when uthread blocks in kernel? Q: do user-level uthreads run concurrently?

big picture:

Multiple threads executing in the kernel Sharing memory, devices, and various data structures. Locks to protect invariants One outstanding disk request One scheduler selecting a thread to run

Show context switching pattern in kernel

sequence coordination:

how to arrange for threads to wait for each other to do e.g., wait for disk interrupt to complete e.g., wait for pipe readers to make space in pipe e.g., wait for child to exit e.g., wait for block to use

straw man solution: spin

waste CPU cycles if need to spin for long time

better solution: primitives for coordination

sleep & wakeup (xv6) condition variables (homework) barriers (homework) etc.

sleep&wakeup:

sleep(chan, lock) sleeps on a "channel", an address to name the condition we are sleeping on

wakeup(chan) wakeup wakes up all threads sleeping on chan this may wake up more than one thread

threads may need to retest the condition they are waiting for to make sure that they don't proceed, if only one thread can proceed therefore sleep is typically called inside a loop

case study: iderw()

designing and implementing these primitives is difficult

- why does sleep take the ide lock as argument? demo: switch acquire(ptable.lock) and release(lk) ie. sleep releases thread's lock before acquiring ptable.lock ideintr() runs before sleeper sets its thread state to SLEEPING it scans proctable but no thread is SLEEPING now sleep acquires ptable.lock and sets current thread to SLEEPING -> sleep misses wakeup; deadlock
 - signals

problem: lost wakeup

wakeup happens before sleeper goes to sleep wakeup is lost many solutions to this problem e.g., count wakeups (as semaphores do) all require some new semantics

xv6 solution: sleep takes a lock as argument

sleeper and wakeup acquires locks for shared data structure sleep holds the lock until after it has ptable.lock once it has ptable.lock, no wakeup can come in before it sets state to sleeping -> no lost wakeup problem requires that sleep takes a lock argument!

iderw example

first hold ptable lock set SLEEP then release the lock argument wakeup cannot get lock until sleeper is already to at sleep why a loop around sleep?

Many primitives in literature to solve lost-wakeup problem

counting wakeup&sleep calls in semaphores pass locks as an extra argument in condition variables (as in sleep) etc.

Another example: pipe

what is the race if sleep didn't take p->lock as argument?

kill: how to kill a process?

problem: target process may be running show you cannot clean it up yet solution: target commits suicide source sets flag target thread checks flag in trap and exits downside: it may take a while until process is really killed e.g., until timer interrupt goes off another complication: what if target thread is sleeping in the kernel

signals and sleep

goal: user wants to kill a process (ctrl-C), but process is at sleep in the kernel hard to get right: a process could sleep somewhere deep in the kernel a signal forces it out of sleep but when out of sleep, it is *not* because the condition it is waiting on is true

xv6 solution

some sleep loops check for p->killed (e.g., see pipe.c) BUT not always (see ide.c) Q: does iderw check for p->killed in sleep loop? Q: what goes wrong if we would modify xv6 to check for p->killed and return from iderw()? A: we might be in the create() system call, doing several disk writes. if iderw() returns, then may leave the on-disk fs in an inconsistent state e.g., block allocated but it doesn't show up in any inode larger problem: a thread may be sleeping deep in the kernel may break invariants if we just bail out xv6 doesn't allow this kill in iderw(): target thread may not be killed until much later

xv6 spec for kill

target will never execute a user instruction target may be killed but with (long) delay

Potential race between sleep & setting p->killed

race: 1. target may check p->killed 2. source sets p->killed 3. target calls sleep effect: target will be killed after reading from disk may take a long time before that happens or never see consoleread() in console.c target is waiting for input from console, but user never types again BUT race is *not* a violation of xv6 spec for kill

Other "solutions" to dealing with sleep and signals

A common approach is use longjmp (unwind the stack), and retry the system call. but also has corner cases that are difficult to get right

Edit By [MaHua](#)