

## 6.828 2014 Lecture 12: File System

lecture plan: file systems API -> disk layout caching

why are file systems useful? durability across restarts naming and organization sharing among programs and users

why interesting? crash recovery performance API design for sharing security for sharing abstraction is useful: pipes, devices, /proc, /afs, Plan 9 so FS-oriented apps work with many kinds of objects you will implement one for JOS

API example -- UNIX/Posix/Linux/xv6/xc: `fd = open("x/y", -); write(fd, "abc", 3); link("x/y", "x/z"); unlink("x/y");`

high-level choices visible in this API objects: files (vs virtual disk, DB) content: byte array (vs 80-byte records, BTree) naming: human-readable (vs object IDs) organization: name hierarchy synchronization: none (vs locking, versions)

a few implications of the API: fd refers to something that is preserved even if file name changes or if file is deleted while open! a file can have multiple links i.e. occur in multiple directories no one of those occurrences is special so file must have info stored somewhere other than directory thus: FS records file info in an "inode" on disk FS refers to inode with i-number (internal version of FD) inode must have link count (tells us when to free) inode must have count of open FDs inode deallocation deferred until last link and FD are gone

let's talk about xv6

FS software layers system calls name ops | FD ops inodes inode cache log buffer cache ide driver

mechanical disks concentric tracks each track is a sequence of sectors, usually 512 bytes ECC on each sector can only read/write whole sectors thus: sub-sector writes are expensive (read-modify-write)

mechanical disk performance rotating platter -- about 10,000 rpm = 166 rps = 6 ms per rotation arm moves in and out to the right track -- "seek" -- about 10 ms latency: time to read or write a single sector? 13 ms to read or write a random sector this is an eternity for CPU -- 13 million instructions random I/O has throughput of only 40 kilobytes/sec -- laughable throughput: bytes/second for big sequential reads and writes? about 500 sectors per track thus  $500 \times 512 / 0.006 = 43$  megabytes/sec big sequential writes have much higher throughput than random sector I/O! thus file systems have to pay attention to data layout SSDs are faster but share low random write performance

disk blocks most o/s use blocks of multiple sectors, e.g. 4 KB blocks = 8 sectors to reduce book-keeping and seek overheads xv6 uses single-sector blocks for simplicity

on-disk layout xv6 file system on 2nd IDE drive; first has just kernel xv6 treats IDE drive as an array of sectors, hides tracks 0: unused 1: super block (size, ninodes) 2: array of inodes, packed into blocks X: block in-used bitmap (0=free, 1=inuse) Y: file/dir content blocks Z: log for transactions end of disk

"meta-data" everything on disk other than file content super block, i-nodes, bitmap, directory content

on-disk inode type (free, file, directory, device) nlink size addrs[12+1]

direct and indirect blocks

example: how to find file's byte 8000? logical block 15 =  $8000 / 512$  3rd entry in the indirect block

each i-node has an i-number easy to turn i-number into inode inode is 64 bytes long byte address on disk:  $2512 + 64 \times \text{inum}$

directory contents directory much like a file but user can't directly write content is array of dirents dirent: inum 14-byte file name dirent is free if inum is zero

you should view FS as an on-disk data structure [tree: dirs, inodes, blocks] with two allocation pools: inodes and blocks

let's look at xv6 in action focus on disk writes illustrate on-disk data structures via how updated

Q: how does xv6 create a file?

```
rm fs.img
```

```
$ echo > a write 4 ialloc (from create sysfile.c; mark it non-free) write 4 iupdate (from create; initialize nlink &c) write 29 writei (from dirlink fs.c, from create)
```

Q: what's in block 4? look at create() in sysfile.c

Q: why *two* writes to block 4?

Q: what is in block 29?

Q: what if there are concurrent calls to ialloc? will they get the same inode? note bread / write / brelse in ialloc bread locks the block, perhaps waiting, and reads from disk brelse unlocks the block

Q: how does xv6 write data to a file?

```
$ echo x > a write 28 balloc (from bmap, from writei) write 420 bzero write 420 writei (from filewrite file.c) write 4 iupdate (from writei) write 420 writei write 4 iupdate
```

Q: what's in block 28? look at writei call to bmap look at bmap call to balloc

Q: what's in block 420?

Q: why the iupdate? file length and addrs[]

Q: why *two* writei+iupdate?

Q: how does xv6 delete a file?

```
$ rm a write 29 writei (from sys_unlink; directory content) write 4 iupdate (from sys_unlink; link count of file) write 28 bfree (from itrunc, from iput) write 4 iupdate (from itrunc) write 4 iupdate (from iput)
```

Q: what's in block 29? sys\_unlink in sysfile.c

Q: what's in block 4?

Q: what's in block 28? look at iput

Q: why three iupdates?

Let's look at the block cache in bio.c block cache holds just a few recently-used blocks

FS calls bread, which calls bget bget looks to see if block already cached if present and not B\_BUSY, return the block if present and B\_BUSY, wait if not present, re-use an existing buffer Q: why goto loop after sleep()?

Q: what is the block cache replacement policy? prev ... head ... next bget re-uses bcache.head.prev -- the "tail" brelse moves block to bcache.head.next

Q: is that the best replacement policy?

Q: what if lots of processes need to read the disk? who goes first? iderw appends to idequeue list ideintr calls idestart on head of idequeue list so FIFO

Q: is FIFO a good disk scheduling policy? priority to interactive programs? elevator sort?

Q: how fast can an xv6 application read big files? contiguous blocks? blow a rotation -- no prefetch?

Q: why does it make sense to have a double copy of I/O? disk to buffer cache buffer cache to user space can we fix it to get better performance?

Q: how much RAM should we dedicate to disk buffers?

Edit By [MaHua](#)