

2014 Lecture 20: OSeS and networking ==

Assigned reading: [IX: A Protected Dataplane Operating System for High Throughput and Low Latency](#)

Plan: Lab 6 Commodity OS networking IX

Lab 6 -- [see design picture in lab handout] No interrupts Many IPCs Much copying of network packet content Any parallelism? Many kernel/User transitions? Where are packet queues? Many scheduling decisions? Could you have long delays to process a packet? Does this design achieve high throughput?

Commodity OS (simplified) -- [draw picture] NIC generates interrupt Interrupt handler runs in top-half runs in interrupt context performs minimal work Kernel schedules Bottom-half runs performs TCP processing copies data into sockets may send some packets from outgoing socket Kernel schedules user-level process User-level process reads data from socket (system call) User-level process writes data to socket (system call) TCP processing as a side-effect of write Write to NIC

IX Paper -- Very recent paper (published two months ago) Builds on Dune paper from last week IX is another example use of Dune Brings a lot of 6.828 topics together Isolation Exokernel architecture Multicore scalability Virtualization Networking

IX

Goal: high performance High packets rates for short messages Data center apps involve many servers Low latency, predictable In data centers some messages are short Setting up/closing connection fast

Current OSeS: fine-grained resource scheduling Multiple applications share a single core System call and interrupt latency > packet interarrival time --> interrupt coalescing, queuing delay, intermediate buffering, CPU scheduling --> adds latency to packet processing --> buffering and synchronization increase memory and CPU overhead, reduce throughput API introduces sharing

Idea: separate control plane from data plane Control plane: responsible for coarse-grained resource allocation Data plane: network stack and app logic Use Dune for isolation of control plane, data plane, and app

IX Control plane Schedules resources among data planes CPU core is dedicated to data plane Memory in large pages are allocated to data plane NIC queues are assigned to dataplane cores Data plane has direct access to NIC queue

Data plane: a libos specialized for networking Run to completion with adaptive batching Allow use of polling, instead of interrupts Adaptive batching every stage of networking processing system call boundary, network API, NIC queue Only batch under high load Why a maximum batch size? Zero-copy API on receive: packets are mapped read-only into application on send: scatter/gather memory list flow control No interaction between cores RRS flow groups are assigned to a dedicated core, which does all processing Each core has its own memory pool API doesn't incur sharing Elastic threads dedicated to core Run networking processing and app processing code Time limit on app processing code Large pages Why is this important?

Isolation between data plane and control User-level stacks No protection between stack and application Stack may interfere with other stacks IX design Linux kernel (with Dune module) runs in root, ring 0 IX runs as a library OS in non-root, ring 0 Application (with libix) in non-root, ring 3 Implementation: Dune Does app have direct access to NIC? Can app write pages after handing it to IX? Can app write messages buffers received from NIC? Who provides the *physical* addresses to NIC? Is this an isolation problem?

Libix: network API (see figure 1b and table 1) low-level API: load network commands into a shared buffer call `run_io()` ix process shared buffer posts results in shared buffer of event conditions poll NIC libevent

Aside: commutativity rule for multicore scalability Recall RadixVM paper: if map/unmap are on different regions, a scalable implementation exists General form: if two operations commute, then a scalable implementation must exist Example: map and unmap on different regions commute thus a scalable implementation exists (e.g., RadixVM) IX uses this rule for the design of its interface to ensure scalability no shared name space of file descriptors cookies instead of fds in table 1

IX Performance Comparing Linux, mTCP, and IX What is goodput? Is 5.7usec good? What is the limiting factor? Impact of non-root ring 0 and 3? Why much faster than Linux? Why 99% metric?

Edit By [MaHua](#)